

# Conducting and Presenting Multiple Linear Regression Analysis Using R

Jishnu Matra



# What is Linear Regression?

## Simple Linear Regression:

- ▶ Predicts one thing based on another thing
- ▶ Example: Predicting house price based on its size
- ▶ Draws a straight line through your data that best fits the pattern
- ▶ Goal: Find the line that makes the smallest prediction errors

## Multiple Linear Regression:

- ▶ Predicts one thing based on multiple factors
- ▶ Example: Predicting house price using size, number of rooms, age, and location
- ▶ Considers all factors together to make better predictions
- ▶ Helps you understand which factors matter most
- ▶ In R, we use the simple `lm()` function to do this

## What is Lasso Regression?

### Lasso (Least Absolute Shrinkage and Selection Operator):

- ▶ A smarter version of regular regression
- ▶ Automatically decides which factors are important and which aren't
- ▶ Kicks out the unimportant factors completely (sets them to zero)
- ▶ Prevents the model from being too complicated

### Key Differences from Standard lm():

- ▶ **lm()**: Uses ALL factors you give it, even if some are useless
- ▶ **Lasso**: Automatically removes useless factors, keeps only the important ones
- ▶ **Lambda ( )**: A dial that controls how aggressive Lasso is in removing factors
- ▶ **Benefit**: Gives you a cleaner, simpler model that's easier to understand and explain
- ▶ **Best for**: When you have many factors and aren't sure which ones really matter

# When to Use Linear Regression

**Best For:** Continuous outcome variables (numbers on a scale)

- ▶ Sales revenue from advertising spend
- ▶ House prices based on size and location
- ▶ Temperature predictions from weather variables
- ▶ Student test scores from study hours

**Key Question It Answers:** How does one variable respond to changes in others?

**Bottom Line:** Know your data! Linear regression needs a continuous outcome variable to work properly.

# About the Analysis

## Packages Used:

- ▶ `glmnet` - For Lasso regression
- ▶ `lm` - Ordinary least squares regression model
- ▶ `MASS` - Contains Boston dataset

## Boston Housing Dataset:

- ▶ 506 observations
- ▶ 14 variables
- ▶ Target: `medv` (median home value)
- ▶ Predictors: crime rate, room number, age, etc.

## Key Concepts

**Coefficients:** Numerical values that represent the relationship strength between each predictor variable and the target variable. A larger absolute value indicates stronger influence.

**R-squared ( $R^2$ ):** R squared is the model's accuracy score, it ranges from 0 to 1.

**Penalty (L1):** In Lasso, add up all the coefficient values (ignoring positive/negative signs) and "if this total gets too big, we'll add a penalty to the model." This can force some coefficients to become exactly zero, effectively performing feature selection.

**Regularization:** A technique to prevent overfitting. Without it, models can become too complicated, memorizing training data instead of learning real patterns.

**Lambda ( ):** Lamda is a hyper parameter, which can use to control how strict and simple a model can be. If the lamda value is more, then it gives more strict penalties and keeps the model simple, and vice-versa.

## Step 1: Install & Load Packages

```
# Install and load packages
if (!require(glmnet)) {
  install.packages("glmnet")
}
if (!require(ggplot2)) {
  install.packages("ggplot2")
}
if (!require(reshape2)) {
  install.packages("reshape2")
}
library(glmnet)
library(MASS)
library(ggplot2)
library(reshape2)
```

**Purpose:** Set up the required libraries for our analysis

## Step 2: Load & Prepare Data

```
# Load and prepare data  
data(Boston)  
head(Boston, 3)
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptr
1	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	
2	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
3	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	

medv

1 24.0

2 21.6

3 34.7

```
# OR load your own CSV file from computer:  
# my_data <- read.csv("path/to/your/file.csv")  
# Example: my_data <- read.csv("C:/Users/YourName/Documents/  
  
X <- as.matrix(Boston[, -14])  
y <- Boston$medv
```

## Step 3: Fit Linear Model

```
# Fit a normal linear model  
lm_model <- lm(medv ~ ., data = Boston)  
lm_coef <- coef(lm_model)  
lm_r_squared <- summary(lm_model)$r.squared  
  
cat("---- Base lm() Coefficients ---\n")
```

--- Base lm() Coefficients ---

```
print(lm_coef)
```

	(Intercept)	crim	zn	indus
02	3.645949e+01	-1.080114e-01	4.642046e-02	2.055863e-02
01	2.686734e+00			
	nox	rm	age	dis
01	-1.776661e+01	3.809865e+00	6.922246e-04	-1.475567e+00
01				3
	tax	ptratio	black	lstat
01	-1.233459e-02	-9.527472e-01	9.311683e-03	-5.247584e-01

## Step 4: Fit Lasso Model

```
# Fit a lasso model (with cross-validation)
cv_model <- cv.glmnet(X, y, alpha = 1)
best_lambda <- cv_model$lambda.min

cat("Best lambda from cross-validation:", best_lambda, "\n")
```

Best lambda from cross-validation: 0.03373254

```
lasso_model <- glmnet(X, y, alpha = 1, lambda = best_lambda)
```

**Key Point:** Cross-validation automatically selects the optimal lambda value

## Step 5: Extract Lasso Results

```
# Extract Lasso coefficients and predictions
lasso_coef <- coef(lasso_model)
lasso_predictions <- predict(lasso_model, newx = X)
lasso_r_squared <- 1 - (sum((y - lasso_predictions)^2) /
                           sum((y - mean(y))^2))

cat("---- Lasso Coefficients ---\n")
```

--- Lasso Coefficients ---

```
print(lasso_coef)
```

14 x 1 sparse Matrix of class "dgCMatrix"

s0

(Intercept)	34.160482304
crim	-0.096703363
zn	0.040702395
indus	.
chas	2.674470344

## Step 6: Compare Results

```
# R-squared Comparison
cat("---- R-squared Comparison ---\n")

--- R-squared Comparison ---
cat("Base lm() R-squared:", lm_r_squared, "\n")

Base lm() R-squared: 0.7406427
cat("Lasso glmnet R-squared:", lasso_r_squared, "\n")

Lasso glmnet R-squared: 0.740005
cat("\nDifference:", lm_r_squared - lasso_r_squared, "\n")

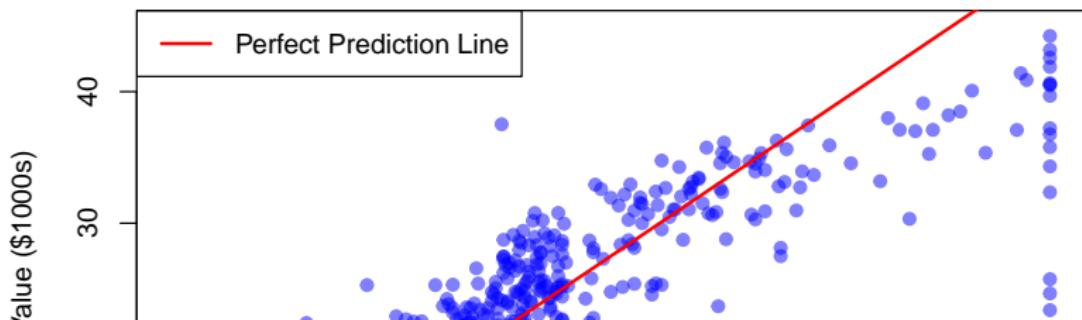
Difference: 0.0006376334
```

**Observation:** Lasso achieves similar performance while potentially simplifying the model by shrinking some coefficients

## Plot: Predicted vs Actual

```
# Predicted vs Actual values
plot(y, lasso_predictions,
      main = "Lasso: Predicted vs Actual Values",
      xlab = "Actual Median Value ($1000s)",
      ylab = "Predicted Median Value ($1000s)",
      pch = 19, col = rgb(0, 0, 1, 0.5))
abline(0, 1, col = "red", lwd = 2)
legend("topleft", legend = "Perfect Prediction Line",
       col = "red", lwd = 2)
```

Lasso: Predicted vs Actual Values



# Business Applications

## Real-World Uses:

### 1. Sales Forecasting

Predict future sales based on marketing spend, pricing, and seasonality to plan inventory and budgets.

### 2. Customer Churn Prediction

Identify which customer behaviors indicate they're likely to leave, so you can take action early.

### 3. Marketing Channel Optimization

Use Lasso to find which advertising channels (social media, TV, email) actually drive revenue and cut the rest.

### 4. Employee Attrition Modeling

Predict which employees might quit based on workload, satisfaction scores, and performance reviews.

### 5. Demand Forecasting

Estimate product demand using weather, promotions, and foot traffic to optimize stock levels.

## Exercise: Try With Another Dataset

### Task:

Run the same workflow (load → prepare → fit → compare) using a *different dataset* such as `mtcars`, `iris`, or your own CSV.

### Steps:

1. Load a dataset: `df <- mtcars`
2. Pick a target variable: `y <- df$mpg`
3. Create predictors: `X <- as.matrix(df[, -1])`
4. Fit linear & Lasso models (same code as earlier)
5. Compare  $R^2$  and identify important predictors.

### Goal:

See how model behavior changes with a new dataset.

Thank You!

**Questions?**