

Conducting and Presenting Multiple Linear Regression Analysis Using R

Jishnu Matra

About the Analysis

Packages Used:

- `glmnet` - For Lasso regression
- `lm` - Ordinary least squares regression model
- `MASS` - Contains Boston dataset

Boston Housing Dataset:

- 506 observations
 - 14 variables
 - Target: `medv` (median home value)
 - Predictors: crime rate, room number, age, etc.
-

Key Concepts

Coefficients: Numerical values that represent the relationship strength between each predictor variable and the target variable. A larger absolute value indicates stronger influence.

R-squared (R^2): A statistical measure (ranging from 0 to 1) that indicates the proportion of variance in the dependent variable that is predictable from the independent variables. Higher values indicate better model fit.

Penalty (L1): In Lasso regression, the penalty is the sum of absolute values of coefficients. This can force some coefficients to become exactly zero, effectively performing feature selection.

Regularization: A technique to prevent overfitting by adding a penalty term to the loss function, discouraging overly complex models with large coefficients.

Lambda (): The regularization strength parameter. Higher values create stronger penalties, leading to more coefficients being shrunk toward zero.

Step 1: Install & Load Packages

```
# Install and load packages
if (!require(glmnet)) {
  install.packages("glmnet")
}
library(glmnet)
library(MASS)
```

Purpose: Set up the required libraries for our analysis

Step 2: Load & Prepare Data

```
# Load and prepare data
data(Boston)
head(Boston, 3)
```

```
      crim zn indus chas   nox     rm    age     dis   rad tax ptratio black lstat
1 0.00632 18  2.31     0 0.538 6.575 65.2 4.0900     1 296   15.3 396.90 4.98
2 0.02731  0  7.07     0 0.469 6.421 78.9 4.9671     2 242   17.8 396.90 9.14
3 0.02729  0  7.07     0 0.469 7.185 61.1 4.9671     2 242   17.8 392.83 4.03
  medv
1 24.0
2 21.6
3 34.7
```

```
X <- as.matrix(Boston[, -14])
y <- Boston$medv
```

What we did: Loaded the Boston dataset and separated predictors (X) from the target variable (y)

Step 3: Fit Linear Model

```
# Fit a normal linear model
lm_model <- lm(medv ~ ., data = Boston)
lm_coef <- coef(lm_model)
lm_r_squared <- summary(lm_model)$r.squared

cat("--- Base lm() Coefficients ---\n")

--- Base lm() Coefficients ---

print(lm_coef)

(Intercept)          crim           zn           indus          chas
3.645949e+01 -1.080114e-01  4.642046e-02  2.055863e-02  2.686734e+00
      nox            rm           age           dis           rad
-1.776661e+01  3.809865e+00  6.922246e-04 -1.475567e+00  3.060495e-01
      tax            ptratio         black         lstat
-1.233459e-02 -9.527472e-01  9.311683e-03 -5.247584e-01

cat("\nR-squared:", lm_r_squared, "\n")
```

R-squared: 0.7406427

Step 4: Fit Lasso Model

```
# Fit a lasso model (with cross-validation)
cv_model <- cv.glmnet(X, y, alpha = 1)
best_lambda <- cv_model$lambda.min

cat("Best lambda from cross-validation:", best_lambda, "\n")
```

Best lambda from cross-validation: 0.007613501

```
lasso_model <- glmnet(X, y, alpha = 1, lambda = best_lambda)
```

Key Point: Cross-validation automatically selects the optimal lambda value

Step 5: Extract Lasso Results

```
# Extract Lasso coefficients and predictions
lasso_coef <- coef(lasso_model)
lasso_predictions <- predict(lasso_model, newx = X)
lasso_r_squared <- 1 - (sum((y - lasso_predictions)^2) /
                           sum((y - mean(y))^2))

cat("--- Lasso Coefficients ---\n")
```

--- Lasso Coefficients ---

```
print(lasso_coef)
```

```
14 x 1 sparse Matrix of class "dgCMatrix"
  s0
(Intercept) 35.841545884
crim        -0.105275607
zn          0.044824451
indus       0.008841511
chas        2.699117520
nox         -17.262113725
rm          3.828134920
age         .
dis         -1.460765357
rad          0.288031707
tax         -0.011418008
ptratio     -0.944235679
black       0.009225021
lstat      -0.522848702
```

Step 6: Compare Results

```
# R-squared Comparison  
cat("--- R-squared Comparison ---\n")
```

```
--- R-squared Comparison ---
```

```
cat("Base lm() R-squared:", lm_r_squared, "\n")
```

```
Base lm() R-squared: 0.7406427
```

```
cat("Lasso glmnet R-squared:", lasso_r_squared, "\n")
```

```
Lasso glmnet R-squared: 0.7405855
```

```
cat("\nDifference:", lm_r_squared - lasso_r_squared, "\n")
```

```
Difference: 5.720643e-05
```

Observation: Lasso achieves similar performance while potentially simplifying the model by shrinking some coefficients

Business Applications

1. Sales Forecasting

Use regression to predict future sales using variables like marketing spend, pricing, or seasonality.

2. Customer Churn Prediction

Identify which customer behaviors or attributes most strongly indicate churn risk.

3. Marketing Channel Optimization

Use Lasso to find which marketing channels contribute most to revenue and remove low-impact ones.

4. Employee Attrition & Performance Modeling

Predict employee outcomes based on workload, satisfaction, and manager ratings.

5. Demand Forecasting for Inventory

Estimate product demand using promotions, weather, or store traffic to improve stock planning.

Exercise: Try With Another Dataset

Task:

Run the same workflow (load → prepare → fit → compare) using a *different dataset* such as `mtcars`, `iris`, or your own CSV.

Steps:

1. Load a dataset: `df <- mtcars`
2. Pick a target variable: `y <- df$mpg`
3. Create predictors: `X <- as.matrix(df[, -1])`
4. Fit linear & Lasso models (same code as earlier)
5. Compare R² and identify important predictors.

Goal:

See how model behavior changes with a new dataset.

Conclusion

- **Linear Regression:** Uses all predictors, no regularization
 - **Lasso Regression:** Applies L1 penalty, can eliminate irrelevant features
 - **Performance:** Both models show similar R² values on this dataset
 - **Advantage of Lasso:** Provides feature selection and helps prevent overfitting
 - **Use Case:** Lasso is particularly useful when you have many predictors and want a simpler, more interpretable model
-

Thank You!

Questions?